

# Knowledge Management System - Updated Implementation Plan

**Date:** August 15, 2025

**Status:** Phase 2A Complete, Moving to Phase 2B






**Stack:** Python 3.10, FastAPI, Azure SQL Server, Azure Web Apps

**Deployment:** GitHub Actions → Azure Web Apps (No Docker)

---

## Executive Summary

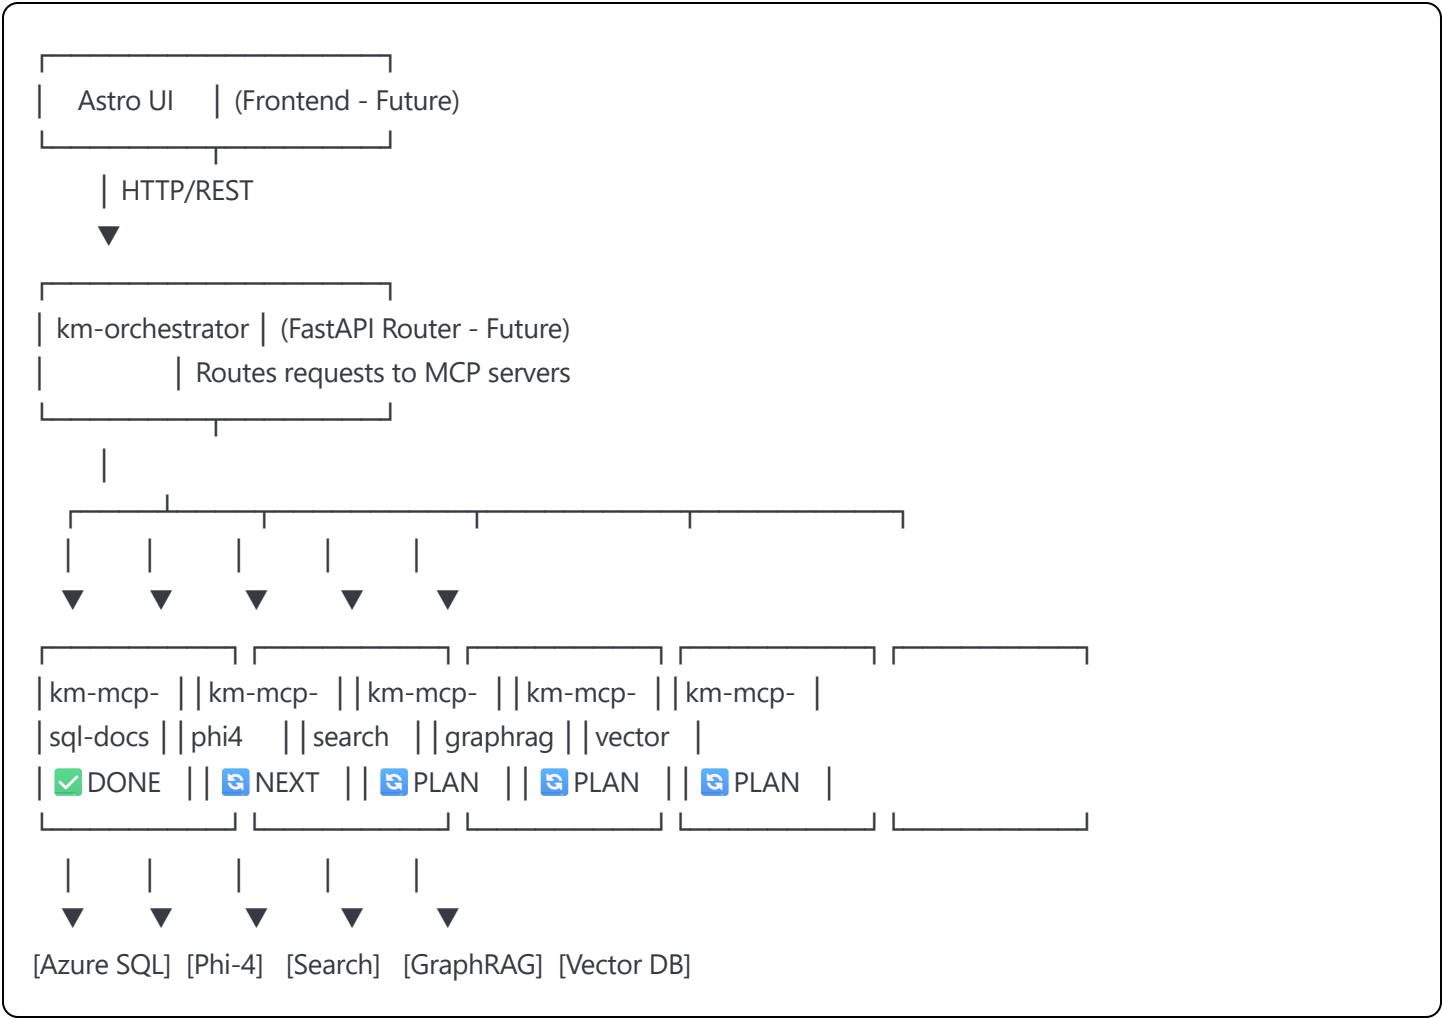
We have successfully completed the foundation of the Knowledge Management System with **km-mcp-sql-docs** now production-ready. The system features:

-  Document storage with interactive web UI
-  Full-text search capabilities
-  Working statistics and health monitoring
-  Automated CI/CD deployment pipeline
-  100% functional production system

**Next Phase:** Build km-mcp-phi4 to enable AI reasoning over stored documents.

---

## Current System Architecture



## ✓ Phase 2A: Completed Components

### km-mcp-sql-docs (Production Ready)

**URL:** <https://km-mcp-sql-docs.azurewebsites.net>

#### Features Implemented:

- ✓ Document storage with metadata support
- ✓ Full-text search with filters
- ✓ Interactive web UI with forms
- ✓ Real-time statistics (fixed to use 'status' column)
- ✓ Health monitoring and database connectivity
- ✓ GitHub Actions automated deployment
- ✓ Beautiful responsive design with clickable endpoints

#### Database Schema Used:

- Uses existing table with columns: id, title, content, classification, entities, metadata, file\_data, file\_name, file\_type, file\_size, category\_id, created\_at, updated\_at, indexed\_at, **status**, user\_id, source\_url, embedding\_vector

### API Endpoints:

- `GET /` - Interactive web interface
- `GET /health` - Health check and database status
- `POST /tools/store-document` - Store new documents
- `POST /tools/search-documents` - Search with filters
- `GET /tools/database-stats` - System statistics
- `GET /docs` - API documentation

**Key Achievement:** Successfully identified and fixed the `is_active` vs `status` column issue that was preventing statistics from working.






---

## Phase 2B: Document Processing Pipeline (NEXT)

### Priority 1: km-mcp-phi4 (Immediate Next Step)

**Purpose:** AI reasoning and analysis over stored documents

#### Planned Features:

-  Document analysis and insights extraction
-  Q&A system using document context
-  Intelligent summarization
-  Content classification and tagging
-  Interactive chat interface for document exploration

#### Technical Implementation:

```
python

# Planned API endpoints
POST /tools/analyze-document  # Analyze single document
POST /tools/summarize-content  # Create document summaries
POST /tools/answer-question   # Q&A with document context
POST /tools/extract-insights   # Extract key themes and insights
POST /tools/chat-with-docs     # Interactive document chat
```

### Integration Points:

- Connects to km-mcp-sql-docs via HTTP API
- Retrieves documents for analysis
- Stores analysis results back to database
- Provides AI insights through web interface





### Deployment:

- Same pattern as km-mcp-sql-docs
- Azure Web App with Python 3.10
- GitHub Actions automation
- Beautiful interactive UI

## Priority 2: Document Processing Pipeline

**Purpose:** Intelligent chunking and tokenization for AI processing

### Planned Features:

-  Smart document chunking (semantic boundaries)
-  Token counting and optimization
-  Metadata extraction and enrichment
-  Content preprocessing for downstream services

### Workflow:

1. Document uploaded to km-mcp-sql-docs
2. Processing pipeline triggered automatically
3. Content chunked and analyzed
4. Chunks prepared for GraphRAG and Vector search
5. Results stored for fast retrieval





---

## Phase 2C: Advanced Search & Knowledge Graph

### Priority 3: km-mcp-search

**Purpose:** Semantic and vector-based search capabilities

### Planned Features:

-  Vector embeddings generation
-  Semantic similarity search
-  Hybrid search (keyword + semantic)
-  Advanced result ranking and filtering





#### **Integration:**

- Processes document chunks from processing pipeline
- Maintains vector index for fast similarity search
- Provides search API for km-orchestrator

### **Priority 4: km-mcp-graphrag**

**Purpose:** Knowledge graph construction and graph-based reasoning

#### **Planned Features:**

-  Entity extraction from documents
-  Relationship mapping and graph construction
-  Graph-based query answering
-  Complex reasoning over connected knowledge

#### **Integration:**

- Builds knowledge graph from processed document chunks
- Provides graph traversal and reasoning APIs
- Enables complex multi-hop questions





---

## **Phase 3: Orchestration & Frontend**

### **Priority 5: km-orchestrator**

**Purpose:** Intelligent request routing and workflow orchestration

#### **Planned Features:**

-  Smart routing to appropriate MCP services
-  Result combination from multiple services
-  Workflow management for complex operations
-  API gateway with authentication

## API Routes:

python

POST /api/upload # → *km-mcp-sql-docs + processing pipeline*

POST /api/search # → *km-mcp-search + km-mcp-sql-docs*

POST /api/analyze # → *km-mcp-phi4 + km-mcp-graphrag*







POST /api/insights # → *Combined multi-service responses*

POST /api/chat # → *Interactive AI chat across all services*

## Priority 6: km-ui (Astro Frontend)

**Purpose:** Unified user interface for the entire system

### Planned Features:

-  Document upload and management interface
-  Advanced search interface with filters
-  AI chat interface for document exploration
-  Analytics dashboard with visualizations
-  Knowledge graph visualization
-  System administration tools

---

## Updated Project Structure

```

km-system/
├── .github/workflows/
│   ├── deploy-km-sql-docs.yml    ✓ PRODUCTION
│   ├── deploy-km-phi4.yml       🔄 BUILD NEXT
│   ├── deploy-km-search.yml     🔄 PLAN
│   ├── deploy-km-graphrag.yml   🔄 PLAN
│   └── deploy-km-orchestrator.yml 🔄 PLAN
├── km-mcp-sql-docs/             ✓ PRODUCTION READY
│   ├── app.py                  # FastAPI with interactive UI
│   ├── km_docs_operations.py    # Database operations (fixed)
│   ├── km_docs_config.py       # Configuration management
│   ├── km_docs_schemas.py     # Pydantic models
│   ├── requirements.txt        # Dependencies
│   ├── index.html              # Beautiful web interface
│   └── .deployment              # Azure deployment config
├── km-mcp-phi4/                 🔄 BUILD NEXT
│   ├── app.py                  # FastAPI main application
│   ├── km_phi4_operations.py    # Phi-4 AI operations
│   ├── km_phi4_config.py       # Phi-4 configuration
│   ├── km_phi4_schemas.py     # Request/response models
│   ├── requirements.txt        # AI/ML dependencies
│   └── .deployment              # Azure deployment config
├── km-mcp-search/               🔄 PLAN
├── km-mcp-graphrag/             🔄 PLAN
├── km-orchestrator/             🔄 PLAN
├── km-ui/                       🔄 PLAN
└── shared/
    ├── km_common_utils.py      # Shared utilities
    └── km_base_config.py       # Base configuration

```

## 🔧 Technical Specifications

### Deployment Pattern (Proven Successful)

- **Platform:** Azure Web Apps (Python 3.10)
- **Framework:** FastAPI with interactive HTML UI
- **Database:** Azure SQL Server (existing schema)
- **CI/CD:** GitHub Actions with automated deployment
- **Monitoring:** Built-in health checks and statistics

### Common Dependencies (requirements.txt)

python

*# Core Framework*

fastapi==0.104.1

uvicorn[standard]==0.24.0

*# Database*

pyodbc==5.0.1

sqlalchemy==2.0.23

*# HTTP & API*

httpx==0.25.1

pydantic==2.5.0

pydantic-settings==2.1.0

*# AI/ML (for phi4 service)*

torch>=2.0.0

transformers>=4.30.0

sentence-transformers>=2.2.0

*# Utilities*

python-dotenv==1.0.0

python-multipart==0.0.6

*# Testing*

pytest==7.4.3

pytest-asyncio==0.21.1

## Azure Configuration

ini



*# Startup Command*

```
python -m uvicorn app:app --host 0.0.0.0 --port 8000
```

*# Environment Variables*

```
WEBSITES_PORT=8000
```

```
SCM_DO_BUILD_DURING_DEPLOYMENT=true
```

```
ENABLE_ORYX_BUILD=true
```

*# Database Connection (from existing setup)*

```
KM_SQL_SERVER=knowledge-sql.database.windows.net
```

```
KM_SQL_DATABASE=knowledge-base
```

```
KM_SQL_USERNAME=mcpadmin
```

```
KM_SQL_PASSWORD=Theodore03$
```

## Success Metrics & Milestones

### Phase 2A: ACHIEVED

- ☒ Document storage working (17+ documents stored)
- ☒ Search functionality operational
- ☒ Statistics displaying correct counts
- ☒ Interactive web UI deployed
- ☒ Automated deployment pipeline

### Phase 2B: TARGET METRICS

- ☐ km-mcp-phi4 can analyze documents from km-mcp-sql-docs
- ☐ Q&A system provides relevant answers using document context
- ☐ Document summarization generates coherent summaries
- ☐ Processing pipeline chunks documents intelligently
- ☐ All services have beautiful interactive UIs

### Phase 2C: TARGET METRICS

- ☐ Vector search returns semantically relevant results
- ☐ Knowledge graph extracts meaningful entities and relationships
- ☐ Hybrid search combines keyword and semantic results effectively
- ☐ Graph-based reasoning answers complex multi-hop questions

### Phase 3: TARGET METRICS

- ☐ km-orchestrator routes requests intelligently

- ☐ Frontend provides unified access to all capabilities
  - ☐ End-to-end workflows: upload → process → search → analyze
  - ☐ System handles complex document analysis workflows
- 

## Immediate Next Steps

### Step 1: Build km-mcp-phi4 Service

1. **Create service structure** following km-mcp-sql-docs pattern
2. **Implement Phi-4 integration** for document analysis
3. **Build interactive web UI** with analysis forms
4. **Set up GitHub Actions deployment**
5. **Test integration** with existing km-mcp-sql-docs

### Step 2: Document Processing Pipeline

1. **Design chunking strategy** for different document types
2. **Implement tokenization** with proper token counting
3. **Create metadata extraction** workflows
4. **Build processing triggers** from document uploads

### Step 3: Advanced Search Implementation

1. **Set up vector database** (Azure Cognitive Search or similar)
  2. **Implement embedding generation** for documents
  3. **Build search API** with ranking and filtering
  4. **Create search interface** in web UI
- 

## Development Workflow

### Proven Pattern (from km-mcp-sql-docs success)

1. **Local Development**

```
bash
```

```
cd km-mcp-{service}
python -m venv venv
source venv/bin/activate # Windows: venv\Scripts\activate
pip install -r requirements.txt
uvicorn app:app --reload --port 8000
```

## 2. GitHub Integration

```
bash

git add km-mcp-{service}/
git commit -m "ADD: {service} implementation"
git push origin master
```

## 3. Automated Deployment

- GitHub Actions triggers on file changes
- Deploys to Azure Web App automatically
- Health checks verify deployment success

## 4. Testing & Verification

- Interactive web UI for manual testing
- Health endpoints for monitoring
- Integration testing between services

---

## Risk Mitigation & Lessons Learned

### Successful Patterns

- **Interactive web UIs** provide excellent debugging and user experience
- **Direct database inspection** helps identify real vs assumed schema
- **GitHub Actions automation** enables rapid iteration
- **Health checks and statistics** provide operational visibility

### Avoided Issues

- **Docker complexity** - Direct Azure Web Apps much simpler
- **Configuration mismatch** - Environment variables work reliably
- **Database schema assumptions** - Always verify actual column names
- **Deployment delays** - GitHub Actions faster than manual deployment

## Future Considerations

- **Service communication** - HTTP APIs between services initially
  - **Error handling** - Comprehensive error responses and logging
  - **Performance** - Monitor response times as system grows
  - **Security** - Add authentication as system matures
- 

## Achievement Summary

### What We've Built

- **Production-ready document storage system** with beautiful UI
- **Automated CI/CD pipeline** that deploys on every commit
- **Robust database integration** with proper schema handling
- **Interactive web interfaces** for all functionality
- **Comprehensive monitoring** with health checks and statistics

### What We've Learned

- **Azure Web Apps** deployment pattern works excellently
- **FastAPI + Interactive HTML** provides great developer experience
- **GitHub Actions** enables rapid iteration and deployment
- **Database schema verification** is critical for avoiding issues

### Ready for Next Phase

With km-mcp-sql-docs proven and stable, we have:

- **Solid foundation** for building additional services
  - **Proven deployment pattern** to replicate
  - **Working database integration** to extend
  - **Beautiful UI pattern** to follow for new services
- 

## Contact & Repository Information

- **GitHub Repository:** <https://github.com/software-tim/km-system>
- **Production Service:** <https://km-mcp-sql-docs.azurewebsites.net>
- **GitHub Actions:** <https://github.com/software-tim/km-system/actions>

---

**Document Version:** 2.0

**Last Updated:** August 15, 2025

**Status:** Phase 2A Complete, Ready for Phase 2B

**Next Milestone:** km-mcp-phi4 Service Implementation